

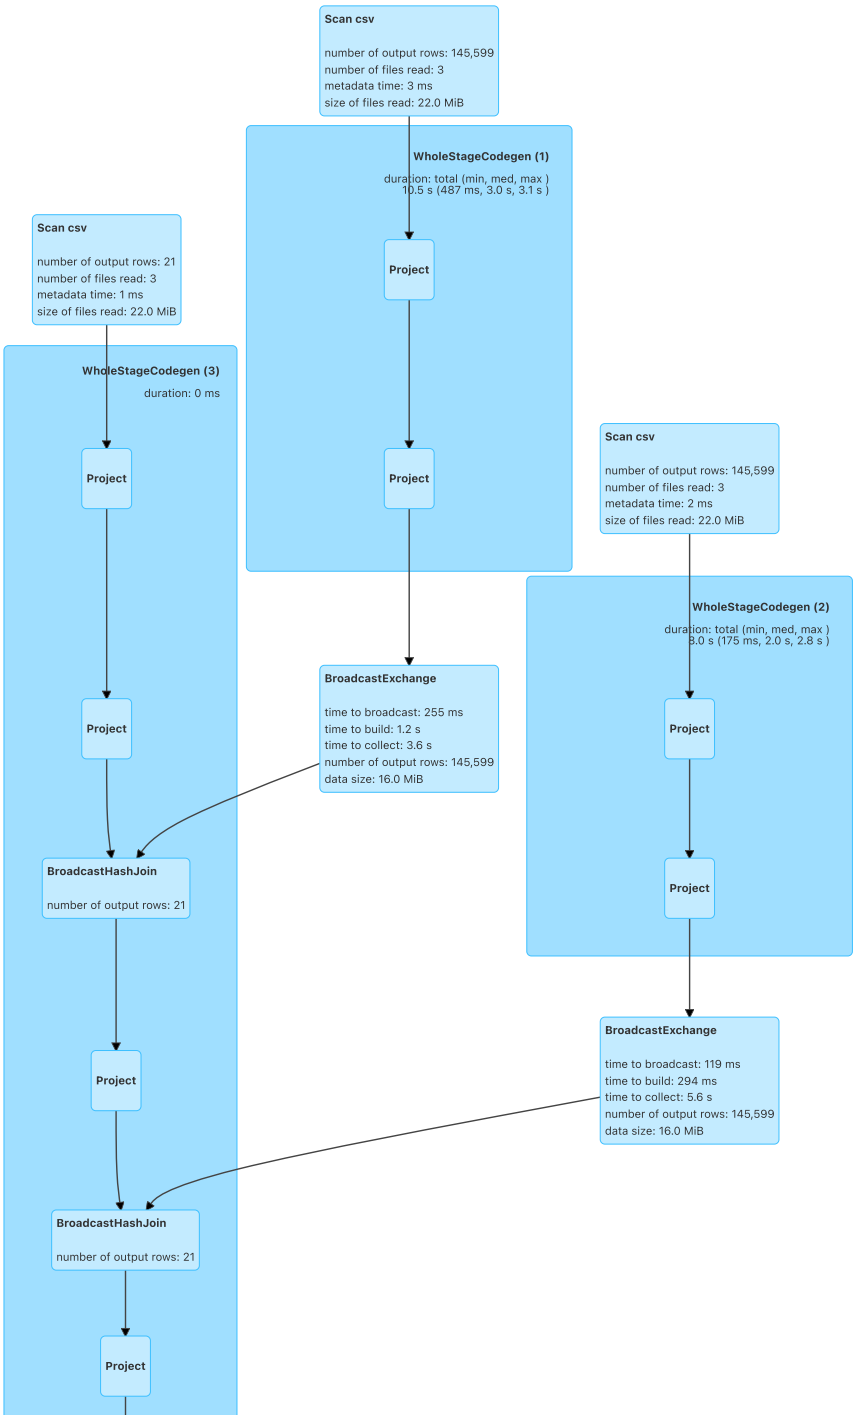
Details for Query 1


Submitted Time: 2024/04/29 16:11:21

Duration: 7 s

Succeeded Jobs: 123

☐ Show the Stage ID and Task ID that corresponds to the max metric




[▼ Details](#)

```

== Physical Plan ==
AdaptiveSparkPlan (32)
+- == Final Plan ==
   CollectLimit (18)
   +- * Project (17)
      +- * BroadcastHashJoin Inner BuildRight (16)
         :- * Project (10)
         :  +- * BroadcastHashJoin Inner BuildRight (9)
         :  :- * Project (3)
         :  :  +- * Project (2)
         :  :  :  +- Scan csv (1)
         :  :  +- BroadcastQueryStage (8), Statistics(sizeInBytes=16.0 MiB, rowCount=1.46E+5)
         :  :  +- BroadcastExchange (7)
         :  :  +- * Project (6)
         :  :  +- * Project (5)
         :  :  +- Scan csv (4)
         +- BroadcastQueryStage (15), Statistics(sizeInBytes=16.0 MiB, rowCount=1.46E+5)
         +- BroadcastExchange (14)
         +- * Project (13)
         +- * Project (12)
         +- Scan csv (11)
+- == Initial Plan ==
   CollectLimit (31)
   +- Project (30)
      +- BroadcastHashJoin Inner BuildRight (29)
         :- Project (25)
         :  +- BroadcastHashJoin Inner BuildRight (24)
         :  :- Project (20)
         :  :  +- Project (19)
         :  :  :  +- Scan csv (1)
         :  :  +- BroadcastExchange (23)
         :  :  +- Project (22)
         :  :  +- Project (21)
         :  :  :  +- Scan csv (4)
         +- BroadcastExchange (28)
         +- Project (27)
         +- Project (26)
         +- Scan csv (11)

(1) Scan csv
Output [3]: [VendorID#17, tpep_pickup_datetime#18, trip_distance#21]
Batched: false
Location: InMemoryFileIndex [file:/content/sample_data/uber_data7.csv, ... 2 entries]
ReadSchema: struct<VendorID:string,tpep_pickup_datetime:string,trip_distance:string>

(2) Project [codegen id : 3]
Output [4]: [VendorID#17, tpep_pickup_datetime#18, trip_distance#21, (monotonically_increasing_id() + 1) AS trip_id#55L]
Input [3]: [VendorID#17, tpep_pickup_datetime#18, trip_distance#21]

(3) Project [codegen id : 3]
Output [6]: [trip_id#55L, cast(VendorID#17 as int) AS VendorID#77, cast(tpep_pickup_datetime#18 as timestamp) AS tpep_pickup_datetime#119, cast(trip_distance#21 as double) AS trip_distance#182, trip_id#55L AS pickup_location_id#478L, trip_id#55L AS dropoff_location_id#455L]
Input [4]: [VendorID#17, tpep_pickup_datetime#18, trip_distance#21, trip_id#55L]

(4) Scan csv
Output [2]: [pickup_longitude#579, pickup_latitude#580]
Batched: false
Location: InMemoryFileIndex [file:/content/sample_data/uber_data7.csv, ... 2 entries]
ReadSchema: struct<pickup_longitude:string,pickup_latitude:string>

(5) Project [codegen id : 1]
Output [3]: [pickup_longitude#579, pickup_latitude#580, (monotonically_increasing_id() + 1) AS trip_id#55L]
Input [2]: [pickup_longitude#579, pickup_latitude#580]

(6) Project [codegen id : 1]
Output [3]: [trip_id#55L AS pickup_location_id#593L, cast(pickup_longitude#579 as double) AS pickup_longitude#203, cast(pickup_latitude#580 as double) AS pickup_latitude#224]
Input [3]: [pickup_longitude#579, pickup_latitude#580, trip_id#55L]

(7) BroadcastExchange
Input [3]: [pickup_location_id#593L, pickup_longitude#203, pickup_latitude#224]
Arguments: HashedRelationBroadcastMode(List(input[0, bigint, false]),false), [plan_id=86]

(8) BroadcastQueryStage
Output [3]: [pickup_location_id#593L, pickup_longitude#203, pickup_latitude#224]
Arguments: 0

(9) BroadcastHashJoin [codegen id : 3]
Left keys [1]: [pickup_location_id#478L]
Right keys [1]: [pickup_location_id#593L]
Join type: Inner
Join condition: None

```

```
(10) Project [codegen id : 3]
Output [8]: [trip_id#55L, VendorID#77, tpep_pickup_datetime#119, trip_distance#182, pickup_location_id#478L, dropff_location_id#455L,
pickup_longitude#203, pickup_latitude#224]
Input [9]: [trip_id#55L, VendorID#77, tpep_pickup_datetime#119, trip_distance#182, pickup_location_id#478L, dropff_location_id#455L,
pickup_location_id#593L, pickup_longitude#203, pickup_latitude#224]

(11) Scan csv
Output [2]: [dropoff_longitude#644, dropoff_latitude#645]
Batched: false
Location: InMemoryFileIndex [file:/content/sample_data/uber_data7.csv, ... 2 entries]
ReadSchema: struct<dropoff_longitude:string,dropoff_latitude:string>

(12) Project [codegen id : 2]
Output [3]: [dropoff_longitude#644, dropoff_latitude#645, (monotonically_increasing_id() + 1) AS trip_id#55L]
Input [2]: [dropoff_longitude#644, dropoff_latitude#645]

(13) Project [codegen id : 2]
Output [3]: [trip_id#55L AS dropff_location_id#654L, cast(dropoff_longitude#644 as double) AS dropoff_longitude#245,
cast(dropoff_latitude#645 as double) AS dropoff_latitude#266]
Input [3]: [dropoff_longitude#644, dropoff_latitude#645, trip_id#55L]

(14) BroadcastExchange
Input [3]: [dropff_location_id#654L, dropoff_longitude#245, dropoff_latitude#266]
Arguments: HashedRelationBroadcastMode(List(input[0, bigint, false]),false), [plan_id=107]

(15) BroadcastQueryStage
Output [3]: [dropff_location_id#654L, dropoff_longitude#245, dropoff_latitude#266]
Arguments: 1

(16) BroadcastHashJoin [codegen id : 3]
Left keys [1]: [dropff_location_id#455L]
Right keys [1]: [dropff_location_id#654L]
Join type: Inner
Join condition: None

(17) Project [codegen id : 3]
Output [11]: [toprettystring(trip_id#55L, Some(Etc/UTC)) AS toprettystring(trip_id)#724, toprettystring(vendorid#77, Some(Etc/UTC)) AS
toprettystring(vendorid)#725, toprettystring(trip_distance#182, Some(Etc/UTC)) AS toprettystring(trip_distance)#726,
toprettystring(tpep_pickup_datetime#119, Some(Etc/UTC)) AS toprettystring(tpep_pickup_datetime)#727,
toprettystring(tpep_pickup_datetime#119, Some(Etc/UTC)) AS toprettystring(tpep_pickup_datetime)#728,
toprettystring(pickup_location_id#478L, Some(Etc/UTC)) AS toprettystring(pickup_location_id)#729, toprettystring(pickup_longitude#203,
Some(Etc/UTC)) AS toprettystring(pickup_longitude)#730, toprettystring(pickup_latitude#224, Some(Etc/UTC)) AS
toprettystring(pickup_latitude)#731, toprettystring(dropff_location_id#455L, Some(Etc/UTC)) AS toprettystring(dropff_location_id)#732,
toprettystring(dropoff_longitude#245, Some(Etc/UTC)) AS toprettystring(dropoff_longitude)#733, toprettystring(dropoff_latitude#266,
Some(Etc/UTC)) AS toprettystring(dropoff_latitude)#734]
Input [11]: [trip_id#55L, VendorID#77, tpep_pickup_datetime#119, trip_distance#182, pickup_location_id#478L, dropff_location_id#455L,
pickup_longitude#203, pickup_latitude#224, dropff_location_id#654L, dropoff_longitude#245, dropoff_latitude#266]

(18) CollectLimit
Input [11]: [toprettystring(trip_id)#724, toprettystring(vendorid)#725, toprettystring(trip_distance)#726,
toprettystring(tpep_pickup_datetime)#727, toprettystring(tpep_pickup_datetime)#728, toprettystring(pickup_location_id)#729,
toprettystring(pickup_longitude)#730, toprettystring(pickup_latitude)#731, toprettystring(dropff_location_id)#732,
toprettystring(dropoff_longitude)#733, toprettystring(dropoff_latitude)#734]
Arguments: 21

(19) Project
Output [4]: [VendorID#17, tpep_pickup_datetime#18, trip_distance#21, (monotonically_increasing_id() + 1) AS trip_id#55L]
Input [3]: [VendorID#17, tpep_pickup_datetime#18, trip_distance#21]

(20) Project
Output [6]: [trip_id#55L, cast(VendorID#17 as int) AS VendorID#77, cast(tpep_pickup_datetime#18 as timestamp) AS tpep_pickup_datetime#119,
cast(trip_distance#21 as double) AS trip_distance#182, trip_id#55L AS pickup_location_id#478L, trip_id#55L AS dropff_location_id#455L]
Input [4]: [VendorID#17, tpep_pickup_datetime#18, trip_distance#21, trip_id#55L]

(21) Project
Output [3]: [pickup_longitude#579, pickup_latitude#580, (monotonically_increasing_id() + 1) AS trip_id#55L]
Input [2]: [pickup_longitude#579, pickup_latitude#580]

(22) Project
Output [3]: [trip_id#55L AS pickup_location_id#593L, cast(pickup_longitude#579 as double) AS pickup_longitude#203,
cast(pickup_latitude#580 as double) AS pickup_latitude#224]
Input [3]: [pickup_longitude#579, pickup_latitude#580, trip_id#55L]

(23) BroadcastExchange
Input [3]: [pickup_location_id#593L, pickup_longitude#203, pickup_latitude#224]
Arguments: HashedRelationBroadcastMode(List(input[0, bigint, false]),false), [plan_id=68]

(24) BroadcastHashJoin
Left keys [1]: [pickup_location_id#478L]
Right keys [1]: [pickup_location_id#593L]
Join type: Inner
Join condition: None
```

```
(25) Project
Output [8]: [trip_id#55L, VendorID#77, tpep_pickup_datetime#119, trip_distance#182, pickup_location_id#478L, dropff_location_id#455L,
pickup_longitude#203, pickup_latitude#224]
Input [9]: [trip_id#55L, VendorID#77, tpep_pickup_datetime#119, trip_distance#182, pickup_location_id#478L, dropff_location_id#455L,
pickup_location_id#593L, pickup_longitude#203, pickup_latitude#224]

(26) Project
Output [3]: [dropoff_longitude#644, dropoff_latitude#645, (monotonically_increasing_id() + 1) AS trip_id#55L]
Input [2]: [dropoff_longitude#644, dropoff_latitude#645]

(27) Project
Output [3]: [trip_id#55L AS dropff_location_id#654L, cast(dropoff_longitude#644 as double) AS dropoff_longitude#245,
cast(dropoff_latitude#645 as double) AS dropoff_latitude#266]
Input [3]: [dropoff_longitude#644, dropoff_latitude#645, trip_id#55L]

(28) BroadcastExchange
Input [3]: [dropff_location_id#654L, dropoff_longitude#245, dropoff_latitude#266]
Arguments: HashedRelationBroadcastMode(List(input[0, bigint, false]),false), [plan_id=72]

(29) BroadcastHashJoin
Left keys [1]: [dropff_location_id#455L]
Right keys [1]: [dropff_location_id#654L]
Join type: Inner
Join condition: None

(30) Project
Output [11]: [toprettystring(trip_id#55L, Some(Etc/UTC)) AS toprettystring(trip_id)#724, toprettystring(vendorid#77, Some(Etc/UTC)) AS
toprettystring(vendorid)#725, toprettystring(trip_distance#182, Some(Etc/UTC)) AS toprettystring(trip_distance)#726,
toprettystring(tpep_pickup_datetime#119, Some(Etc/UTC)) AS toprettystring(tpep_pickup_datetime)#727,
toprettystring(tpep_pickup_datetime#119, Some(Etc/UTC)) AS toprettystring(tpep_pickup_datetime)#728,
toprettystring(pickup_location_id#478L, Some(Etc/UTC)) AS toprettystring(pickup_location_id)#729, toprettystring(pickup_longitude#203,
Some(Etc/UTC)) AS toprettystring(pickup_longitude)#730, toprettystring(pickup_latitude#224, Some(Etc/UTC)) AS
toprettystring(pickup_latitude)#731, toprettystring(dropff_location_id#455L, Some(Etc/UTC)) AS toprettystring(dropff_location_id)#732,
toprettystring(dropoff_longitude#245, Some(Etc/UTC)) AS toprettystring(dropoff_longitude)#733, toprettystring(dropoff_latitude#266,
Some(Etc/UTC)) AS toprettystring(dropoff_latitude)#734]
Input [11]: [trip_id#55L, VendorID#77, tpep_pickup_datetime#119, trip_distance#182, pickup_location_id#478L, dropff_location_id#455L,
pickup_longitude#203, pickup_latitude#224, dropff_location_id#654L, dropoff_longitude#245, dropoff_latitude#266]

(31) CollectLimit
Input [11]: [toprettystring(trip_id)#724, toprettystring(vendorid)#725, toprettystring(trip_distance)#726,
toprettystring(tpep_pickup_datetime)#727, toprettystring(tpep_pickup_datetime)#728, toprettystring(pickup_location_id)#729,
toprettystring(pickup_longitude)#730, toprettystring(pickup_latitude)#731, toprettystring(dropff_location_id)#732,
toprettystring(dropoff_longitude)#733, toprettystring(dropoff_latitude)#734]
Arguments: 21

(32) AdaptiveSparkPlan
Output [11]: [toprettystring(trip_id)#724, toprettystring(vendorid)#725, toprettystring(trip_distance)#726,
toprettystring(tpep_pickup_datetime)#727, toprettystring(tpep_pickup_datetime)#728, toprettystring(pickup_location_id)#729,
toprettystring(pickup_longitude)#730, toprettystring(pickup_latitude)#731, toprettystring(dropff_location_id)#732,
toprettystring(dropoff_longitude)#733, toprettystring(dropoff_latitude)#734]
Arguments: isFinalPlan=true
```